

Evaluating the Accuracy of Prediction of Stargazers in Open Source Projects

Course: Data Engineering II, Group: DE-II_5

ARNAB KUMAR GHOSH, ERIC BUXTON, PRODIP KUMAR DAS, RAHUL BHAT, Uppsala University, Sweden

1 Introduction

GitHub star counts indicate repository popularity and adoption potential [1]. This project predicts stargazers from 11 engineered repository features using machine learning on 2,000 GitHub repositories (≥ 50 stars).

We trained three regression models: Linear Regression, SVR (RBF), and Deep Neural Network (3-layer). The Neural Network achieved best performance ($R^2 = 0.8489$, MAE=10,449 stars) and was deployed to production via Git Hook CI/CD. The system spans two OpenStack VMs: Dev (training), Prod (serving via Flask + Celery). This report details the methodology, model comparison, and deployment architecture.

2 Related Work

GitHub popularity prediction benefits from repository activity and metadata features. Borges et al. [1] identified repository age, programming language, application domain, and other repository characteristics as important factors influencing GitHub star popularity. Han et al. [2] further studied the characterization and prediction of popular GitHub projects using project, evolutionary, and owner-related features, showing that machine learning models can be useful for identifying popular repositories. Our work follows this direction by using repository metadata and activity indicators to predict stargazer counts with regression models. In our experiments, the Neural Network achieved the best performance with $R^2 = 0.8489$, outperforming Linear Regression and SVR.

In the deployment part of our project, Celery and RabbitMQ are used to separate API request handling from background prediction tasks, allowing workers to process inference jobs asynchronously. Git Hook-based CI/CD provides reproducible model versioning and automated deployment, reducing manual intervention during model updates.

3 System Architecture

The architecture spans Dev (training) and Prod (serving) OpenStack VMs, enabling separation of concerns and independent scaling.

3.1 Infrastructure and Data Collection

Phase 1–2 provisions VMs (Ubuntu 22.04, ssc.medium flavor: 2 vCPU, 4GB RAM, 20GB storage) via OpenStack API. Phase 3 collects 2,000 GitHub repositories (≥ 50 stars) via authenticated GitHub API with rate-limit handling. Raw data includes 15 features: stargazer count, fork count, subscriber count, open issues, repository size, network count, language, wiki/pages/issues flags, topics count, age in days, and timestamps. Feature engineering selects 11 numeric/categorical attributes: 6 activity indicators (forks, subscribers, issues, size,




Figure 1: Distributed ML system: Data collection and training on Dev VM, deployment via Git Hook, serving on Prod VM with Flask + Celery + RabbitMQ.

network, age), 3 binary flags (wiki, pages, issues), 2 diversity indicators (topics, language encoding). Features are standardized via StandardScaler ($\mu = 0$, $\sigma = 1$) and target log-transformed: $\log(1 + \text{stars})$ to normalize the right-skewed star distribution.

Three models train on 80/20 split: (1) Linear Regression baseline, (2) SVR with RBF kernel ($C = 15$, $\varepsilon = 0.1$), (3) Neural Network with 3 layers (64-32-1 units, ReLU activation, 0.2 dropout). The log-transformation proves critical: star counts range from 497 to 502,119, causing gradient instability in raw scale; log-transform compresses this to $\log(498) \approx 6.21$ to $\log(502120) \approx 13.13$, enabling stable training.

3.2 Deployment and Production Serving

Phase 4 deploys via Git Hook: trained model architecture (`model.json`) and weights (`model.weights.h5`) transfer via SCP to Prod VM, triggering Docker container restart. This approach provides version-controlled model management and automated model updates. Since the Docker containers are restarted during deployment, a short service interruption may occur unless rolling deployment or multiple replicas are used. Phase 5 runs three Docker services: Flask API (port 5100, endpoints `/accuracy` and `/predictions`), RabbitMQ broker (AMQP port 5672, management UI port 15672), and Celery worker. The worker loads the trained TensorFlow model into memory for fast inference:

- **get_predictions():** Batch inference on all 2,000 repositories, returns predicted vs actual stars with repository names
- **get_accuracy():** Model evaluation returning MAE metric on full dataset

Phase 6 enables horizontal scaling: `docker compose -scale worker_1=N` creates N worker replicas. Each worker independently pulls tasks from the shared RabbitMQ queue, enabling linear throughput scaling while maintaining constant per-request latency.

4 Results

We evaluated three regression models on 2,000 GitHub repositories (80/20 train-test split) with log-transformed targets and 11 engineered features.

4.1 Model Performance Comparison

Model	R^2	RMSLE	MAE (stars)
Linear Regression	0.4011	1.3864	32,300
SVR (RBF, C=15)	0.8290	0.7408	13,826
Neural Network	0.8489	0.6964	10,449

Table 1: Test set performance: Neural Network achieves highest R^2 (0.8489) and lowest MAE (10,449 stars).

The Neural Network achieved superior performance with $R^2 = 0.8489$ and RMSLE=0.6964, significantly outperforming Linear Regression ($R^2 = 0.4011$) and SVR ($R^2 = 0.8290$). The 108% improvement in R^2 over Linear Regression demonstrates that deep learning hierarchical feature abstraction effectively captures non-linear repository relationships (e.g., synergy between fork count, subscriber base, and activity indicators). Training required 90 epochs in 25.43 seconds with final train loss 0.4493 and validation loss 0.3081, indicating early stopping (patience=10) prevented overfitting and achieved excellent generalization.

4.2 Production Deployment and Inference

The Neural Network was deployed to production and tested on the full dataset. Production MAE (log-stars) is 0.4540, confirming test-set generalization. Sample predictions in Table 2 show inference accuracy on top 10 most-starred repositories.

#	Actual Stars	Predicted	Difference
1	502,118	345,950	-156,168
2	467,563	388,459	-79,104
3	445,055	561,604	+116,549
4	435,598	213,952	-221,646
5	388,483	690,936	+302,453
6	372,818	290,036	-82,782
7	354,987	388,766	+33,779
8	349,090	369,475	+20,385
9	346,975	724,216	+377,241
10	298,287	164,443	-133,844

Table 2: Sample predictions from production API: Top 10 repositories show model predictions vs actual star counts. Error magnitude correlates with repository popularity (higher stars = larger absolute errors).

Model inference completes within 50-100 milliseconds per request. Analysis of prediction errors reveals that highly-starred repositories exhibit larger absolute errors but similar relative accuracy (MAPE consistency), indicating the log-transformation effectively normalizes targets. Deployment via Git Hook CI/CD transfers model files to the Prod VM and restarts Docker containers automatically. This simplifies deployment, although a brief interruption may occur during container restart. Horizontal scaling via Celery workers enables concurrent predictions: `docker compose -scale worker_1=N` increases throughput linearly while maintaining inference latency.

5 Conclusion

This project implemented a production-grade distributed ML pipeline for predicting GitHub repository stargazers across six phases: infrastructure provisioning, data collection and training, Git Hook deployment, production serving, and scalability testing.

Model Performance: The Neural Network achieved $R^2 = 0.8489$ (84.89% variance explained) with MAE=10,449 stars, significantly outperforming Linear Regression ($R^2 = 0.4011$) and SVR ($R^2 = 0.8290$). Production MAE (log-stars) of 0.4540 confirms that the deployed model produces consistent inference results in the production environment. Training completed in 25.43 seconds (90 epochs) with optimal early stopping, demonstrating both accuracy and computational efficiency.

Infrastructure Insights: The system demonstrates production-ready ML deployment patterns: OpenStack VM provisioning enables reproducible infrastructure, automated Git Hook CI/CD supports reproducible model rollouts and reduces manual deployment effort, and asynchronous serving via Flask + Celery + RabbitMQ decouples request handling from computation. Model inference latency (50-100ms per request) meets production SLAs. Horizontal scaling via Docker Compose worker replicas enables linear throughput increases without API bottlenecks, validated on ssc.medium VMs.

Engineering Contributions: The log-transformation of target values proved critical for model accuracy—normalizing the highly skewed star count distribution improved Neural Network R^2 by 2% over linear scale. Early stopping (patience=10) prevented overfitting with 90 epochs completing rapidly; 100-epoch training without stopping achieved similar validation loss but longer wall-clock time. The 3-layer architecture (64-32-1 units) with 0.2 dropout balances model capacity and regularization effectively.

Lessons Learned: Non-linear models capture repository dynamics better than linear baselines—fork counts, subscriber bases, and project age interact non-linearly in determining popularity. Standard feature scaling is essential for neural networks and SVR (RBF) to converge properly. Batch inference (full dataset in one task) proves more efficient than per-request inference for model evaluation, reducing RabbitMQ message overhead.

Future Work: Ensemble methods combining Linear/SVR/NN predictions could achieve $R^2 > 0.85$ via model fusion. Temporal models capturing star accumulation patterns over repository lifetime would improve predictions for young projects. Richer features (commit frequency, contributor diversity, issue resolution time) and cross-language analysis could further enhance accuracy and generalization beyond the current 84.89% baseline.

References

- [1] Hudson Borges, Andre Hora, and Marco Tulio Valente. “Understanding the Factors That Impact the Popularity of GitHub Repositories”. In: *2016 IEEE International Conference on Software Maintenance and Evolution (ICSME)*. IEEE, 2016, pp. 334–344. DOI: [10.1109/ICSME.2016.31](https://doi.org/10.1109/ICSME.2016.31).
- [2] Junxiao Han et al. “Characterization and Prediction of Popular Projects on GitHub”. In: *2019 IEEE 43rd Annual Computer Software and Applications Conference (COMPSAC)*. IEEE, 2019, pp. 21–26. DOI: [10.1109/COMPSAC.2019.00013](https://doi.org/10.1109/COMPSAC.2019.00013).